

ECE 650

Systems Programming & Engineering

Spring 2019

Programming with Network Sockets

Rabih Younes
Duke University

Slides are adapted from Brian Rogers and Tyler Bletsch (Duke)

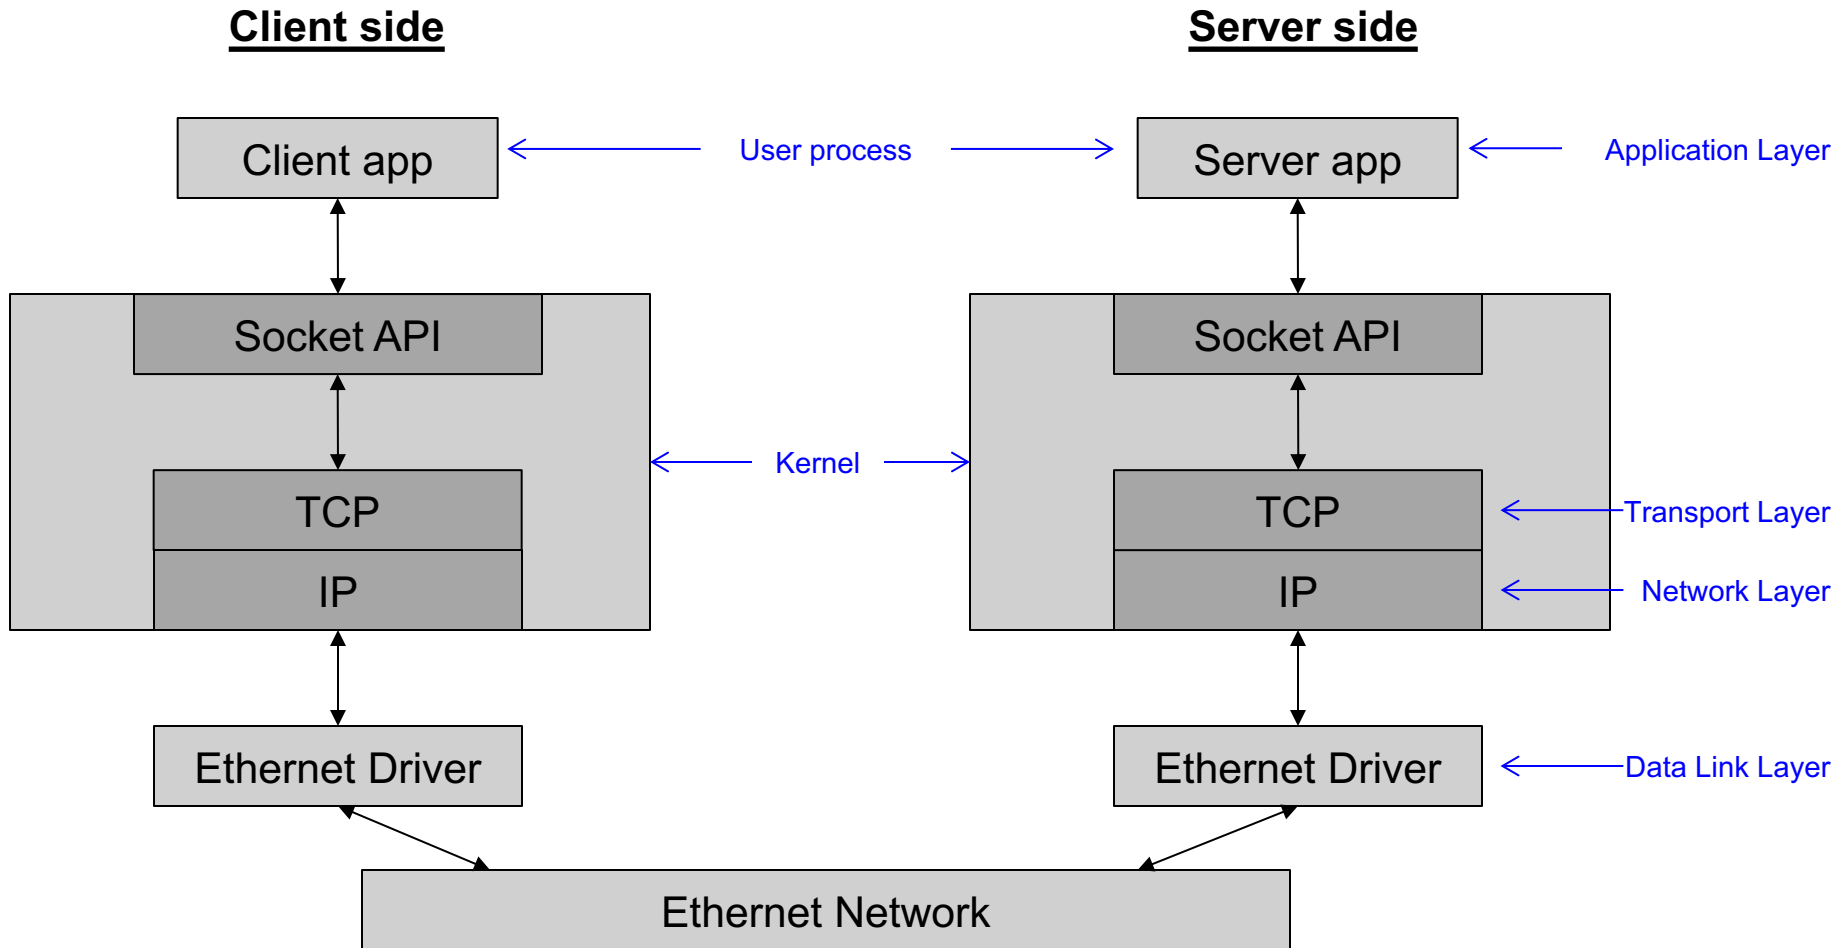
Sockets

- We've looked at shared memory vs. message passing
 - All on a single system (meaning running under a single OS)
- What about communication across distributed processes?
 - Running on different systems
 - Assume systems are connected by a network (e.g. the internet)
- We can program using **network sockets**
 - For creating connections and sending / receiving messages
 - Often follows a client / server pattern
- We will assume basic network knowledge
 - E.g. what is an IP address
 - We will cover the networking stack in more detail in next lectures

Client-Server Model

- Common communication model in networked systems
 - Client typically communicates with a server
 - Server may connect to multiple clients at a time
- Client needs to know:
 - Existence of a server providing the desired service
 - Address (commonly IP address) of the server
- Server does not need to know about the existence of clients and their addresses
 - It waits for service requests from clients
 - It responds using the addresses in the request packets
 - The source address in the client's request would become the destination address in the server's response

Client-Server Overview



- Client and Server communicating across Ethernet using TCP/IP

TCP – Connection-Oriented Service

- **Transmission Control Protocol**

- Designed for end-to-end byte stream **over unreliable network**
- Robust against failures and changing network properties

- TCP transport entity

- E.g., Library procedure(s), user processes, or a part of the kernel
- Manages TCP streams and interfaces (OSI layer 4) to the IP layer (OSI layer 3)
- Accepts user data streams from processes
- **Breaks up data** into pieces that can fit in 1 Ethernet frame (w/ IP + TCP headers)
- Sends each piece separately as IP datagram
- Destination machine TCP entity reconstructs original byte stream
- **Handles retransmissions & re-ordering**

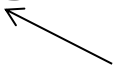
- **Connection-oriented** transport layer

- Provides **error-free, reliable communication**
- Can think of communication between two processes on different machines as just like UNIX pipes or fifos
 - One process puts data in one end, other process takes it out

Network Sockets

- Network interface is identified by an IP address
 - Or a hostname, which translates into an IP address
 - E.g., 127.0.0.1 (same as “localhost”) or login.oit.duke.edu
- Interface has 65536 ports (0-65535)
- Processes attach to ports to use network services
 - Port attachment is done with `bind()` operation
- Allows application-level multiplexing of network services
 - E.g., SSH vs. Web vs. Email use different ports
 - Many ports are standard (e.g., 80 for web server, 22 for SSH)
 - You may have seen URLs like <http://127.0.0.1:4444>
 - 127.0.0.1 is the IP address
 - 4444 is the port number

TCP Service Model

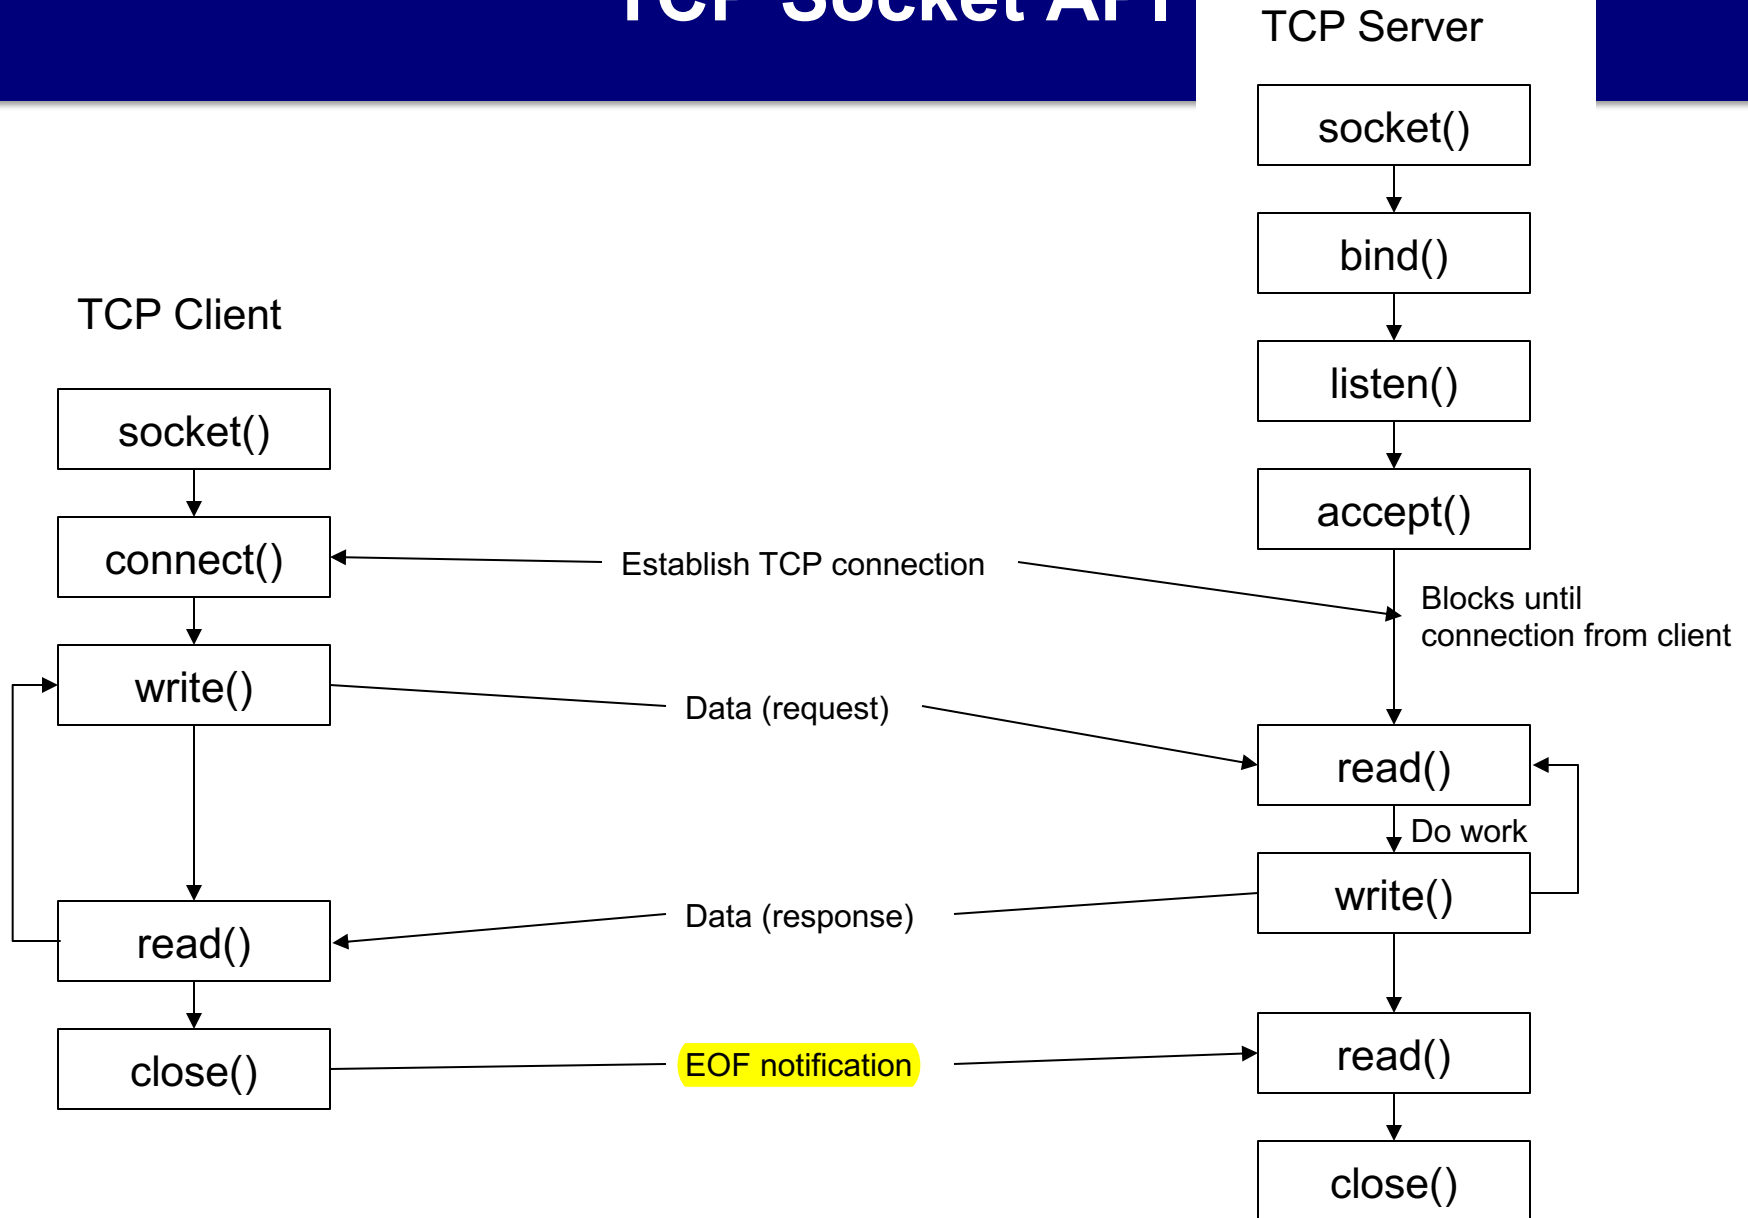
- TCP service setup:
 - Two endpoint processes create endpoints (sockets)
 - Each socket has an address:
 - IP address (32 bit for IPv4 or 128 bits for IPv6)
 - Port number (16 bits)
 - API functions used to create & communicate on sockets
- Port numbers:
 - Numbers below 1024 called “well-known ports”
 - Reserved for standard services, like FTP, HTTP, SMTP
<http://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml>
 - But not all services usually used & active all at once
 - Don't want them all active, just waiting for incoming connections
 - Special daemon: inetd (Internet daemon)  (Still around, but less common nowadays)
 - Attaches to multiple ports
 - Waits for incoming connection
 - fork()'s of the new, appropriate process to handle that connection

UNIX TCP sockets

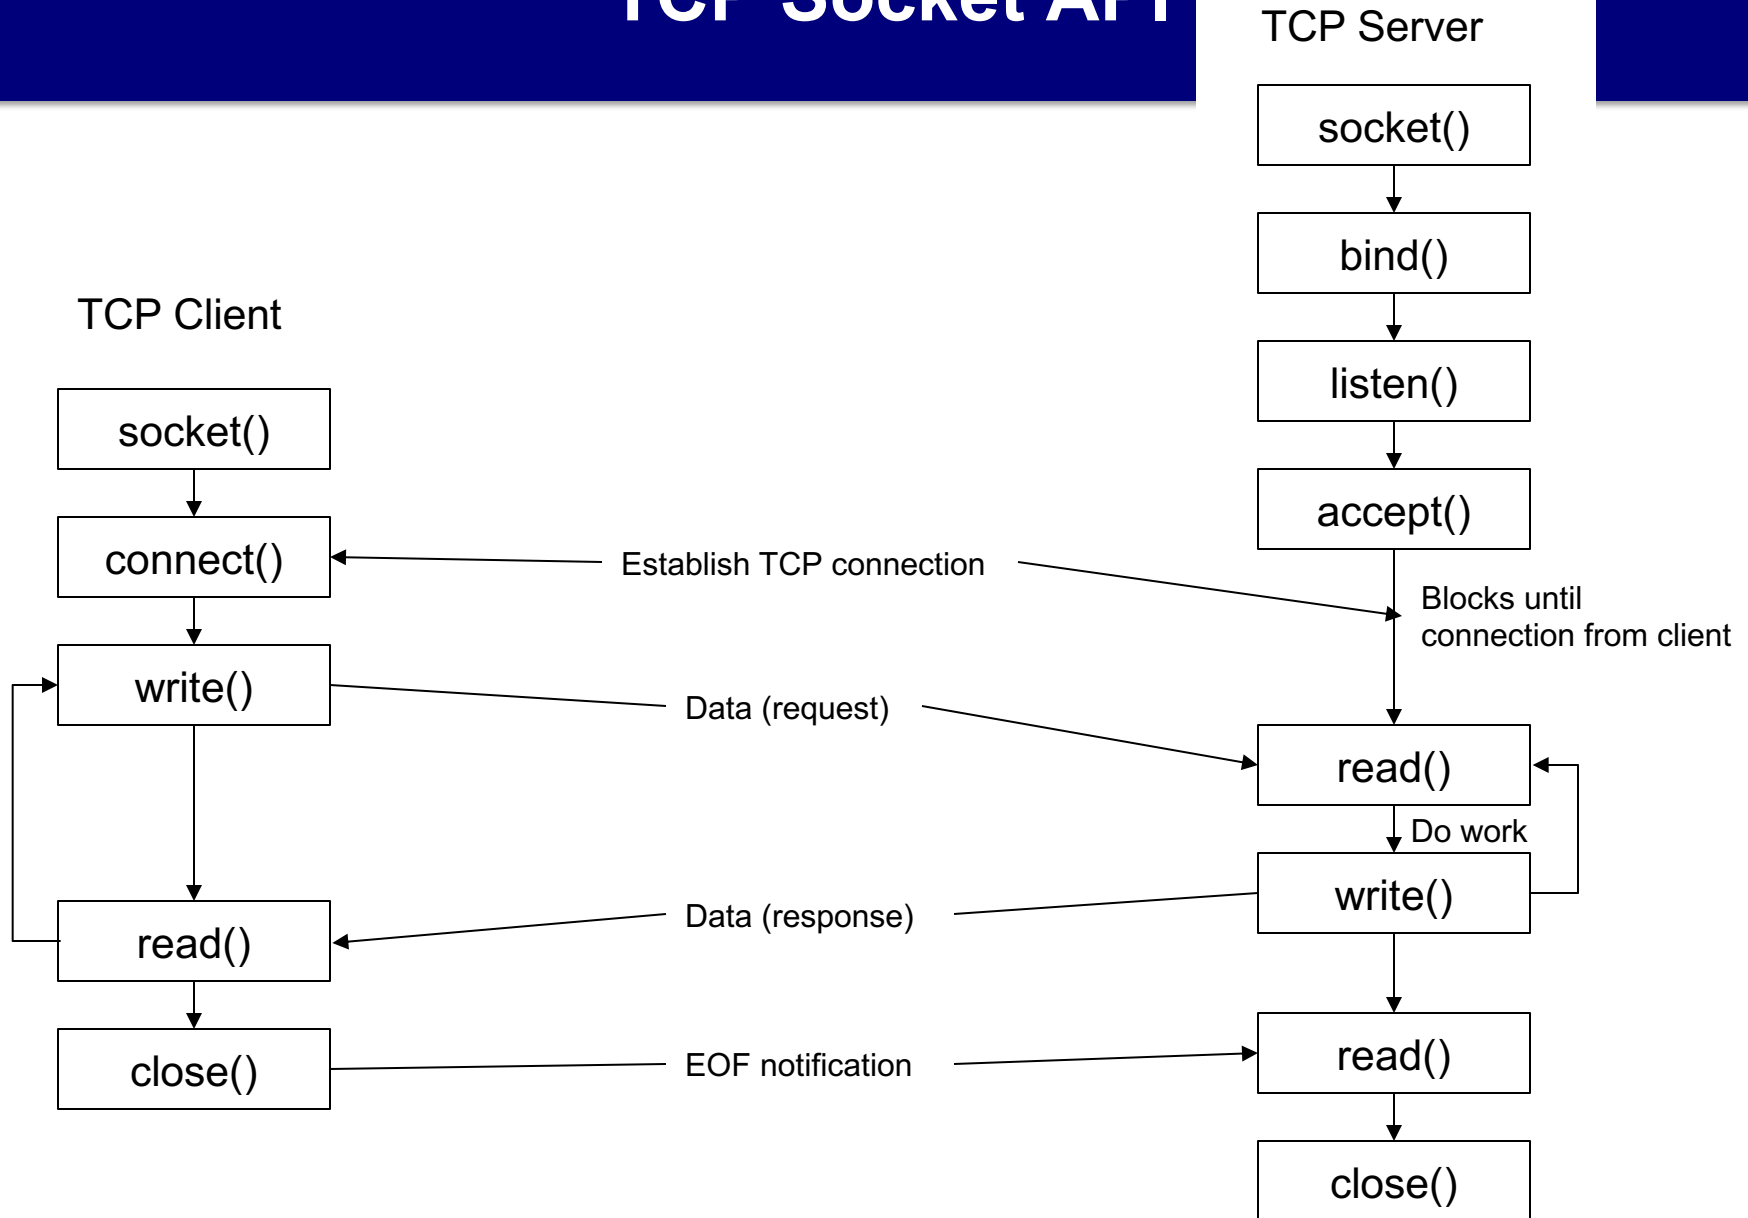
- Here is a great reference for use of socket-related calls
 - <http://beej.us/guide/bgnet/>

Primitive	Meaning
<code>socket()</code>	Create a new communication end point
<code>bind()</code>	Attach a local address to a socket
<code>listen()</code>	Announce willingness to accept connections; give queue size
<code>accept()</code>	Block the caller until a connection attempt arrives
<code>connect()</code>	Actively attempt to establish a connection
<code>send()</code>	Send some data over the connection
<code>recv()</code>	Receive some data from the connection
<code>close()</code>	Release the connection

TCP Socket API



TCP Socket API



Server-Side Structure

- Often follows a common pattern to serve incoming requests

```
pid_t pid;
int listenfd, connfd;
listenfd = socket(...);

/**fill the socket address with server's well known port**/

bind(listenfd, ...);
listen(listenfd, ...);

for ( ; ; ) {

    connfd = accept(listenfd, ...); /* blocking call */

    if ( (pid = fork()) == 0 ) { /* create a child process to service */

        close(listenfd); /* child closes listening socket */

        /**process the request doing something using connfd **/
        /* ..... */

        close(connfd);
        exit(0); /* child terminates
    }
    close(connfd); /*parent closes connected socket*/
}
}
```

Example

- tcp_example
 - server.cpp
 - client.cpp
- Read readme.txt first to know how to execute the files